Started on	Thursday, 24 April 2025, 12:10 PM
State	Finished
Completed on	Thursday, 24 April 2025, 12:15 PM
Time taken	4 mins 51 secs
Marks	14.00/15.00
Grade	**93.33** out of 100.00

Question 1

Complete

Mark 1.00 out of 1.00

Are micro-partitions user-configurable in Snowflake?

- a. No
- b. Yes

Question 2

Complete

Mark 1.00 out of 1.00

How does Snowflake determine which micro-partitions to scan during a query?

- a. Scans all micro-partitions
- b. Uses metadata filters based on pruning
- c. Uses clustering keys
- d. Applies machine learning

Question 3

Complete

Mark 1.00 out of 1.00

How does Snowflake handle changes in data distribution (e.g., skewed data)?

- a. Manual re-partitioning
- b. Auto-reclustering (with clustering keys)
- c. Rewrites old partitions
- d. Requires data export and import

Question 4

Complete

Mark 1.00 out of 1.00

Micro-partitions store data in which format?

- a. Row-based format
- b. JSON
- c. Columnar format
- d. Proprietary Snowflake log

Question 5

Complete

Mark 1.00 out of 1.00

What information does Snowflake store for each micro-partition?

- a. All of the above
- b. Min/Max values per column
- c. Count of NULLs per column
- d. Data skew distribution

Question 6

Complete

Mark 1.00 out of 1.00

What is a Micro-Partition in Snowflake?

- a. A query optimization technique
- b. An automatically created contiguous storage unit
- c. A user-defined partition of data
- d. A block of storage used to store metadata only

Question 7

Complete

Mark 1.00 out of 1.00

What is the advantage of smaller micro-partitions in Snowflake?

- a. Improved write performance
- b. More granular pruning and faster queries
- c. Reduced storage cost
- d. Better support for transactions

Question 8

Complete

Mark 1.00 out of 1.00

What is the typical size range of a Snowflake micro-partition?

- a. 1 MB to 10 MB (compressed)
- b. 1 KB to 5 MB
- c. 10 GB and above
- d. 100 MB to 1 GB

Question 9

Complete

Mark 1.00 out of 1.00

What kind of data structure is used to store metadata about micro-partitions?

- a. B-Trees
- b. Column statistics and ranges
- c. JSON
- d. CSV indexes

Question 10

Complete

Mark 1.00 out of 1.00

When you insert new data into a table, how are micro-partitions affected?

- a. New micro-partitions are automatically created
- b. All data is re-partitioned
- c. Partitions stay unchanged
- d. Existing partitions are overwritten

Question 11

Complete

Mark 1.00 out of 1.00

Which of the following best describes “partition pruning” in Snowflake?

- a. Dropping unused partitions
- b. Skipping micro-partitions that don't match query filters
- c. Rewriting partitions
- d. Caching frequent partitions

Question 12

Complete

Mark 0.00 out of 1.00

Which of the following best describes the immutability of micro-partitions?

- a. They are recreated on each insert
- b. They are deleted after every query
- c. They are read-only after creation
- d. They are mutable but updated in batches

Question 13

Complete

Mark 1.00 out of 1.00

Which of the following can improve the effectiveness of micro-partition pruning?

- a. Using well-designed clustering keys
- b. Querying without WHERE clauses
- c. Using semi-structured data
- d. Writing to the same table continuously

Question 14

Complete

Mark 1.00 out of 1.00

Which of the following tools can help monitor micro-partition behavior in Snowflake?

- a. Query Profiler
- b. Information Schema
- c. Storage Usage Dashboard
- d. SYSTEM$CLUSTERING_INFORMATION function

Question 15

Complete

Mark 1.00 out of 1.00

Which Snowflake feature heavily relies on micro-partition metadata for optimization?

- a. Failover regions
- b. Query Result Caching
- c. Materialized Views
- d. Automatic Clustering